

# Task Parallelism

Prof. Guido Araujo

# Task Parallelism

- Programming model
- Runtime design
- Task stealing
- Runtime performance
- Runtime overhead
- Optimization
  - Task coalescing
  - Dynamic worker deactivation
  - Workload balancing
  - Dependency analysis
- Example

# Programming models

- BSC OpenSs
- OpenMP 4.x
  - Intel IOMP
  - GCC GOMP
  - Unicamp MTSP

```
#pragma omp task in(v[i-1]) out(v[i])  
fun1(&v[i-1], &v[i])
```

# Task Parallelism

## The OpenMP 4.0 task programming model

```
for (int i=1, j=1; i<N; i++) {  
    #pragma omp task in(v[i-1]) out(v[i])  
    fun1(&v[i-1], &v[i]);  
  
    for (int k=0; k<i; k++, j++) {  
        #pragma omp task in(v[i]) out(u[j])  
        fun2(&v[i], &u[j]);  
    }  
  
    fun3( 3 * i );  
}
```


→ Create a task!

→ Create a task!

→ Create NO tasks!

# Task Parallelism


```
for (int i=1, j=1; i<N; i++) {  
    #pragma task in(v) out(v)  
    ● fun1(&v[i-1], &v[i]);  
  
    for (int k=0; k<i; k++, j++) {  
        #pragma task in(v) out(u)  
        ● fun2(&v[i], &u[j]);  
    }  
  
    fun3( 3 * i );  
}
```



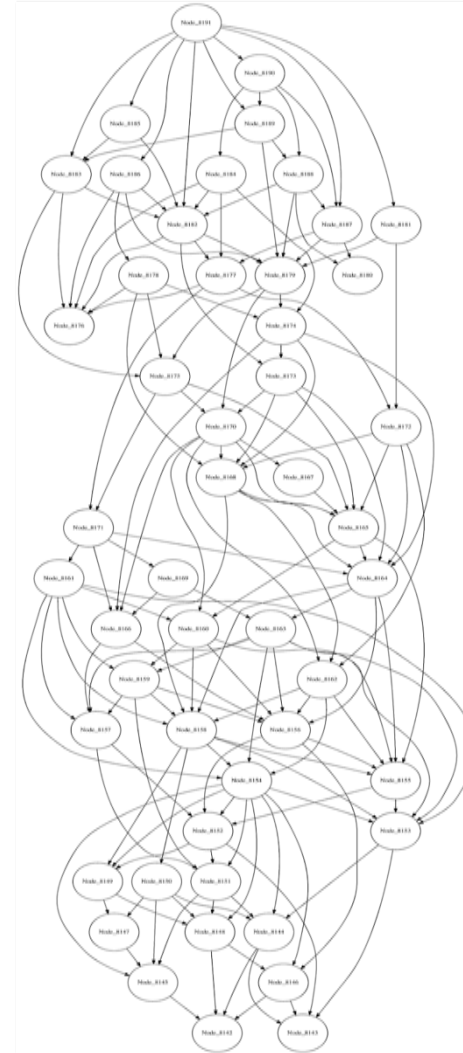
# Task Parallelism Advantages

No need to communicate data. Just detect the dependency


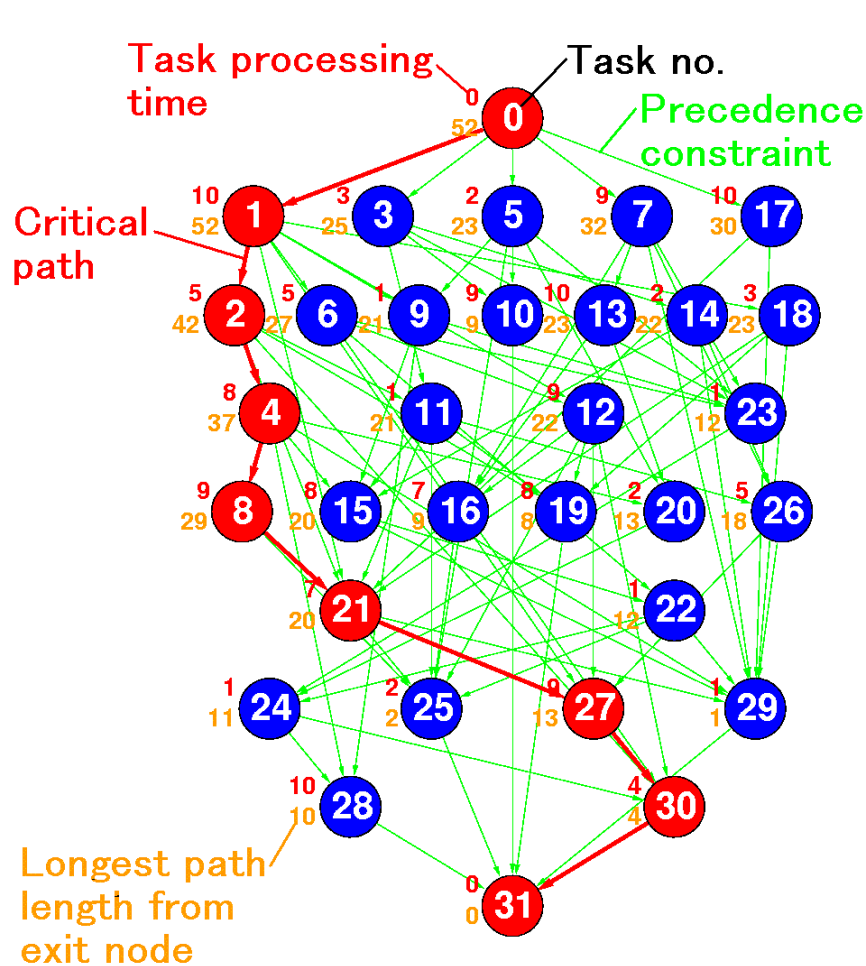
Parallelism between iterations of different loops



# But graphs are not always that simple!



# But graphs are not always that simple!



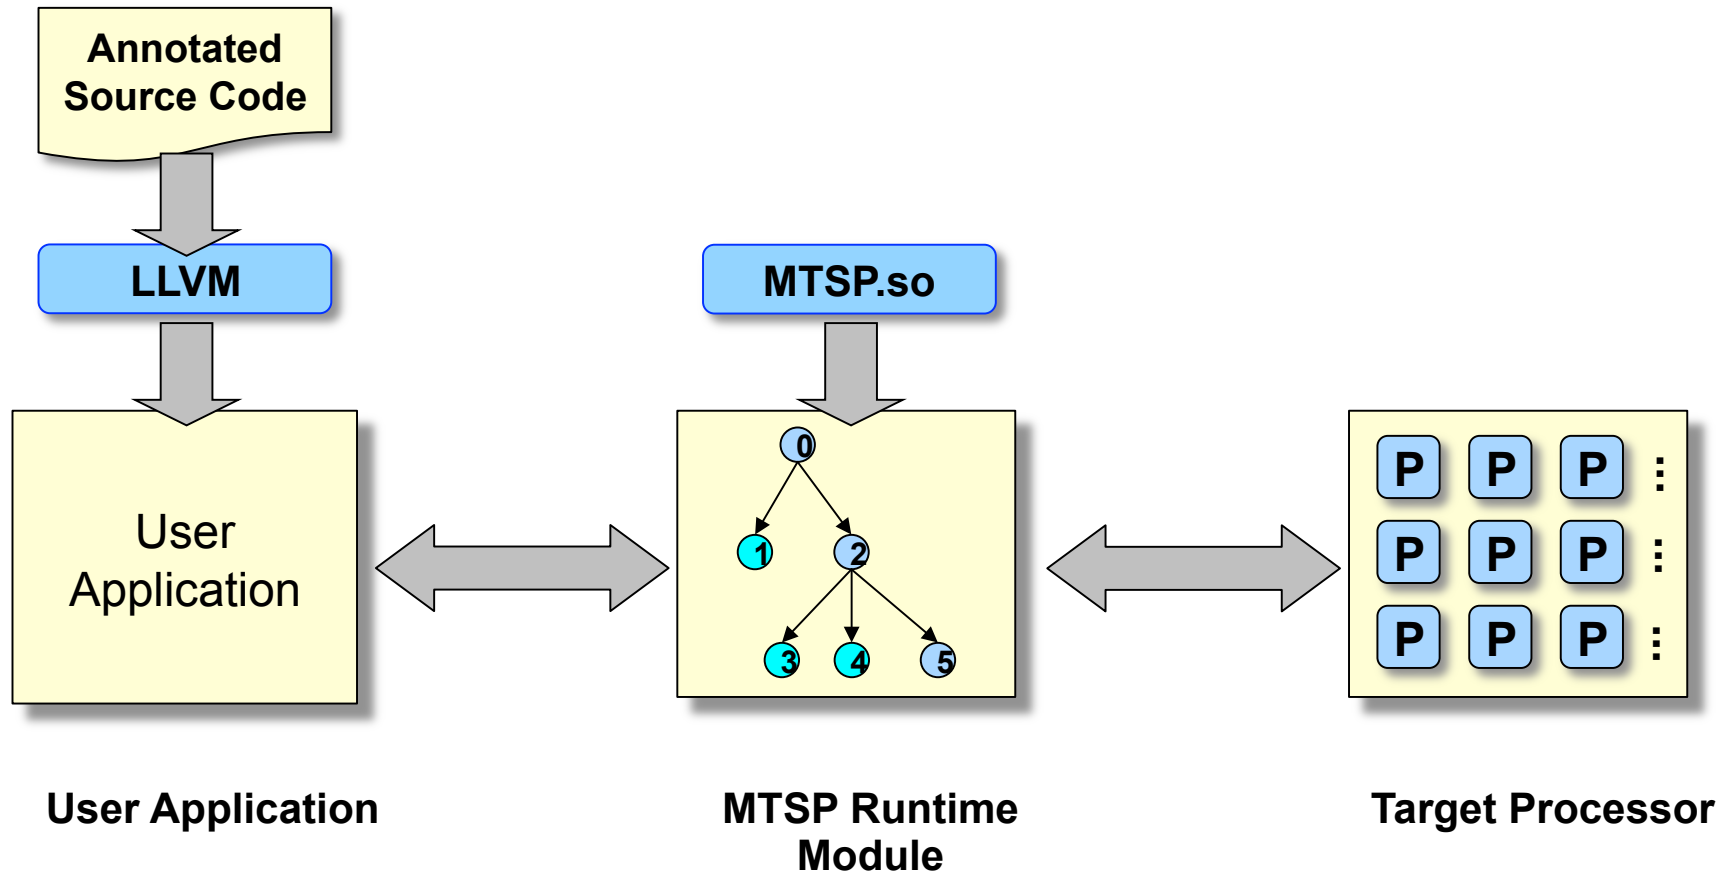


# Task Parallelism


- Programming model
- Runtime design
- Task stealing
- Runtime performance
- Runtime overhead
- Optimization
  - Task coalescing
  - Dynamic worker deactivation
  - Workload balancing
  - Dependency analysis
- Example

# Runtime design

## The MSTP runtime

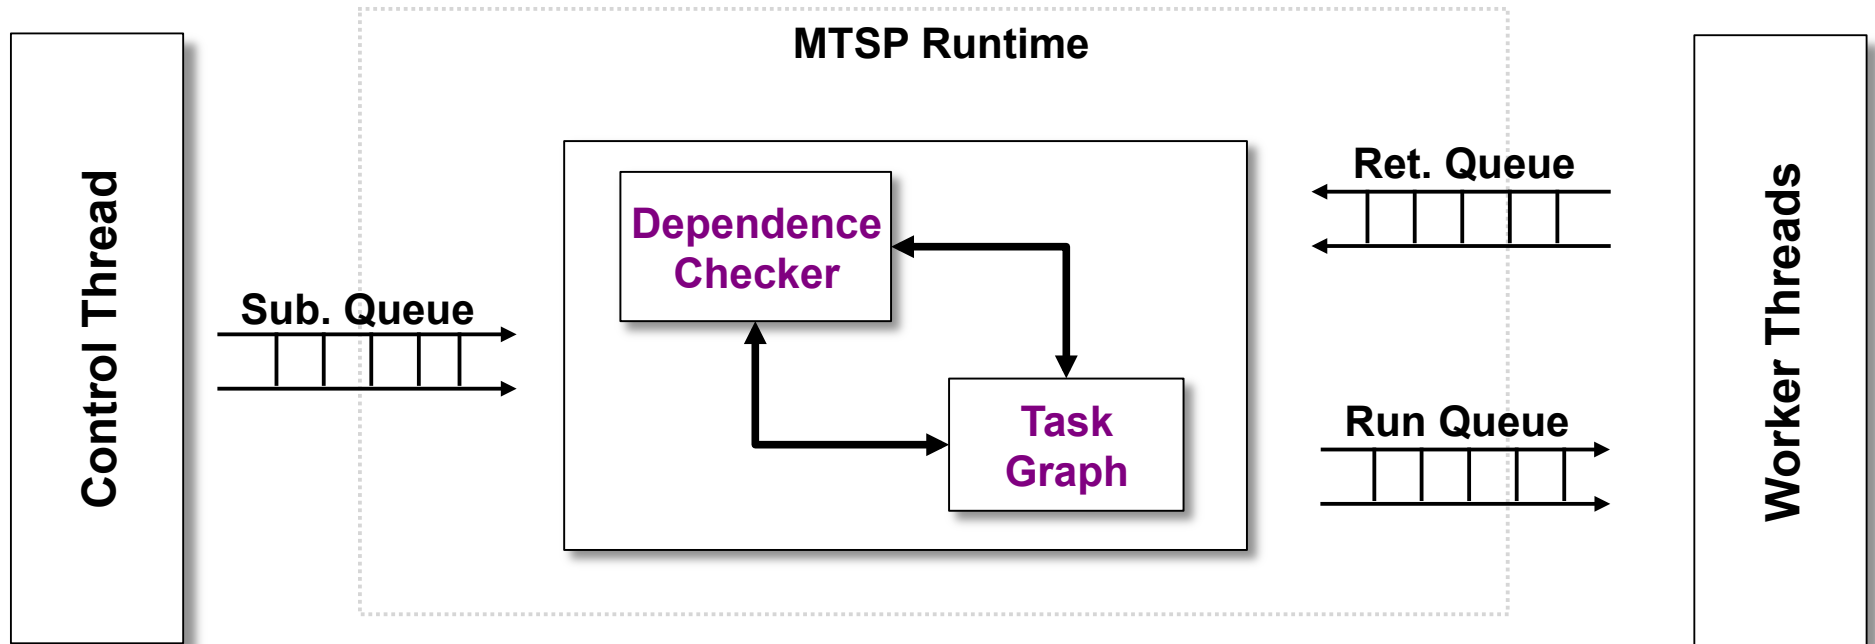


# The CLang OMP Interface

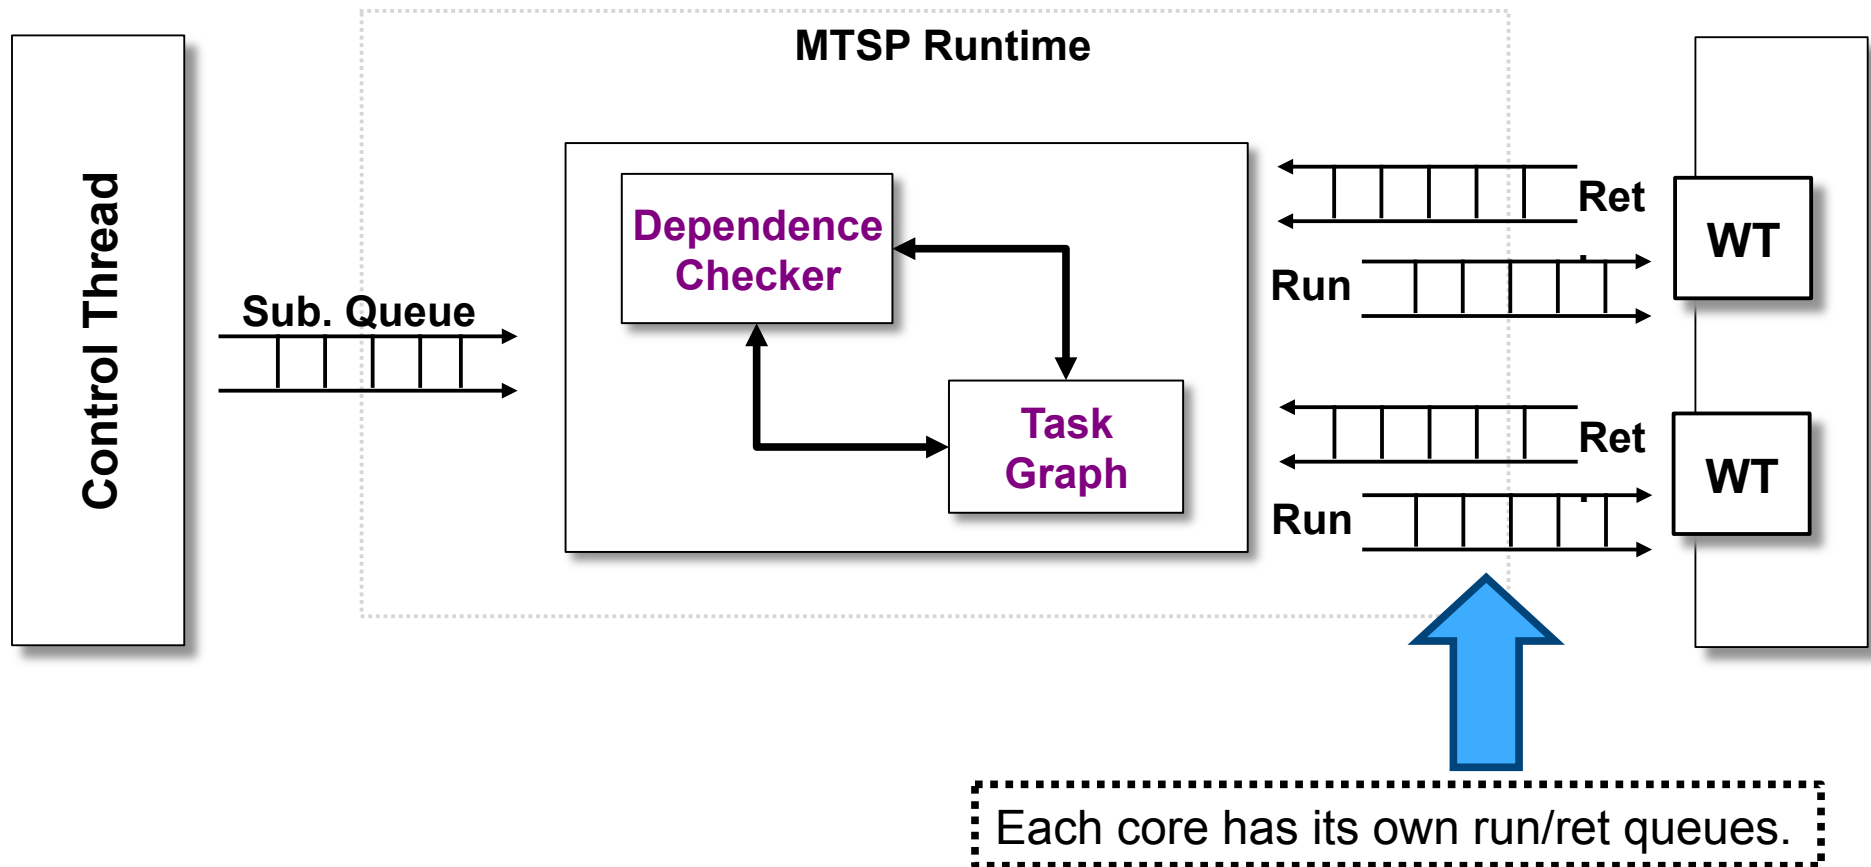


```
/// Used to create a pool of threads  
/// and initiate the system  
void kmpc_fork_call(...);  
/// Used for task allocation  
kmp_task* kmpc_omp_task_alloc(...);  
/// Used to add a task to the runtime  
kmp_int32 kmpc_omp_task_with_deps(...);
```

# MTSP Overview




# Distributed Queues



# Evaluating Workload

Problem: Load Balancing

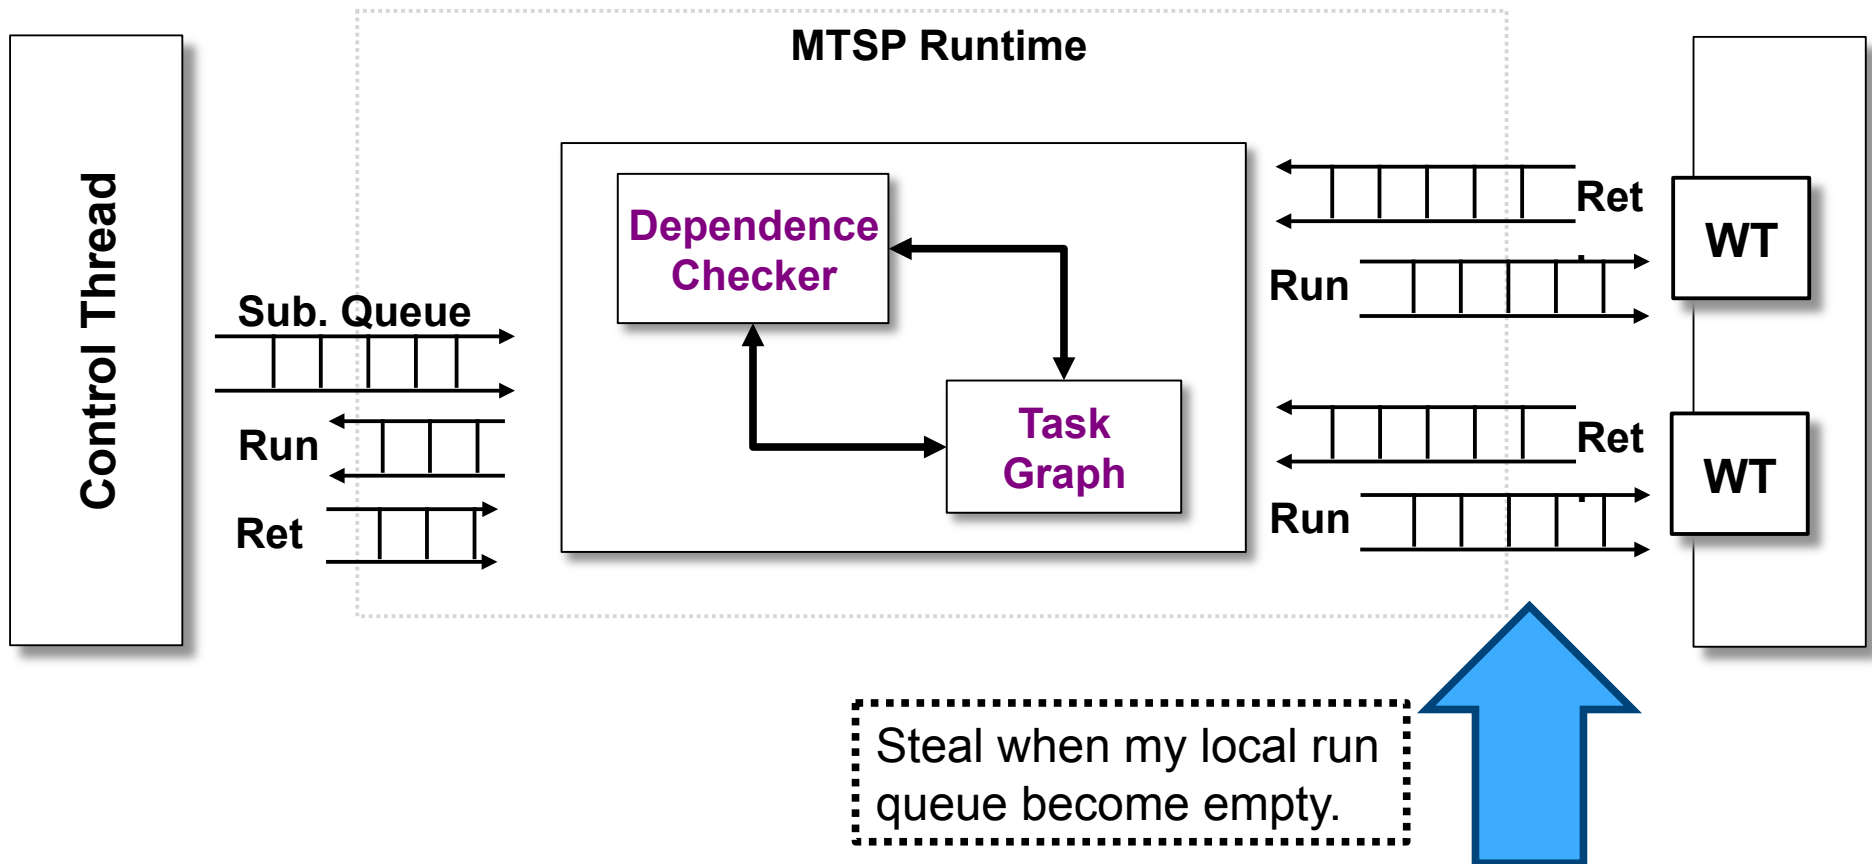


SparseLu -n 64 -m 64

# Task Parallelism

- Programming model
- Runtime design
- Task stealing
- Runtime performance
- Runtime overhead
- Optimization
  - Task coalescing
  - Dynamic worker deactivation
  - Workload balancing
  - Dependency analysis
- Example

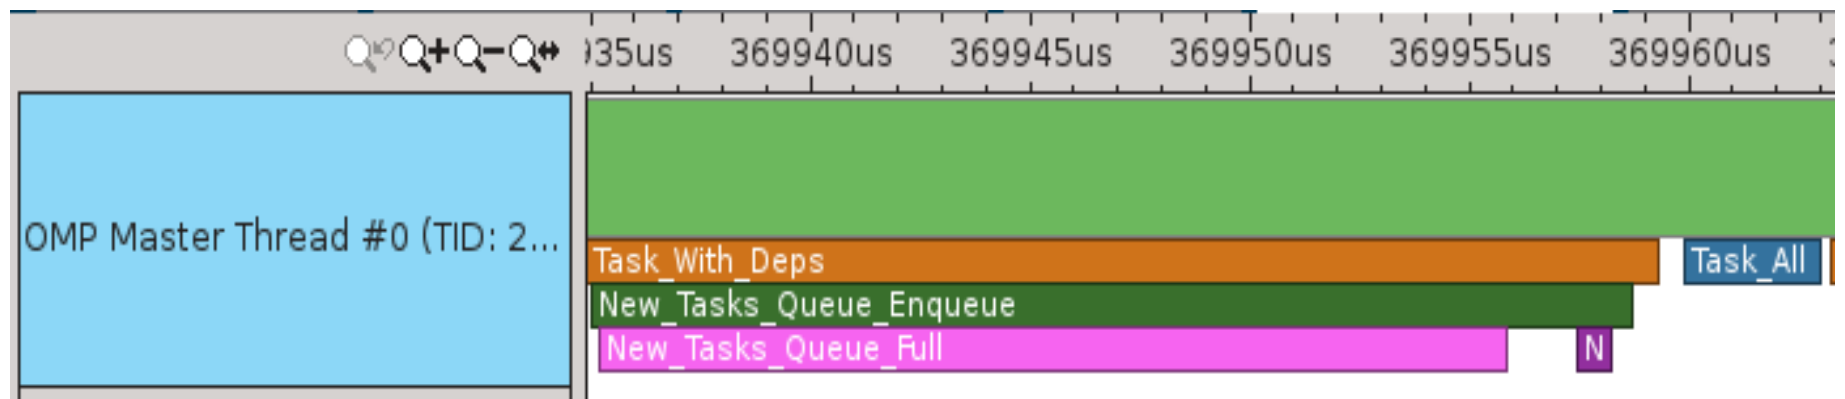

# Work Stealing among Worker Threads



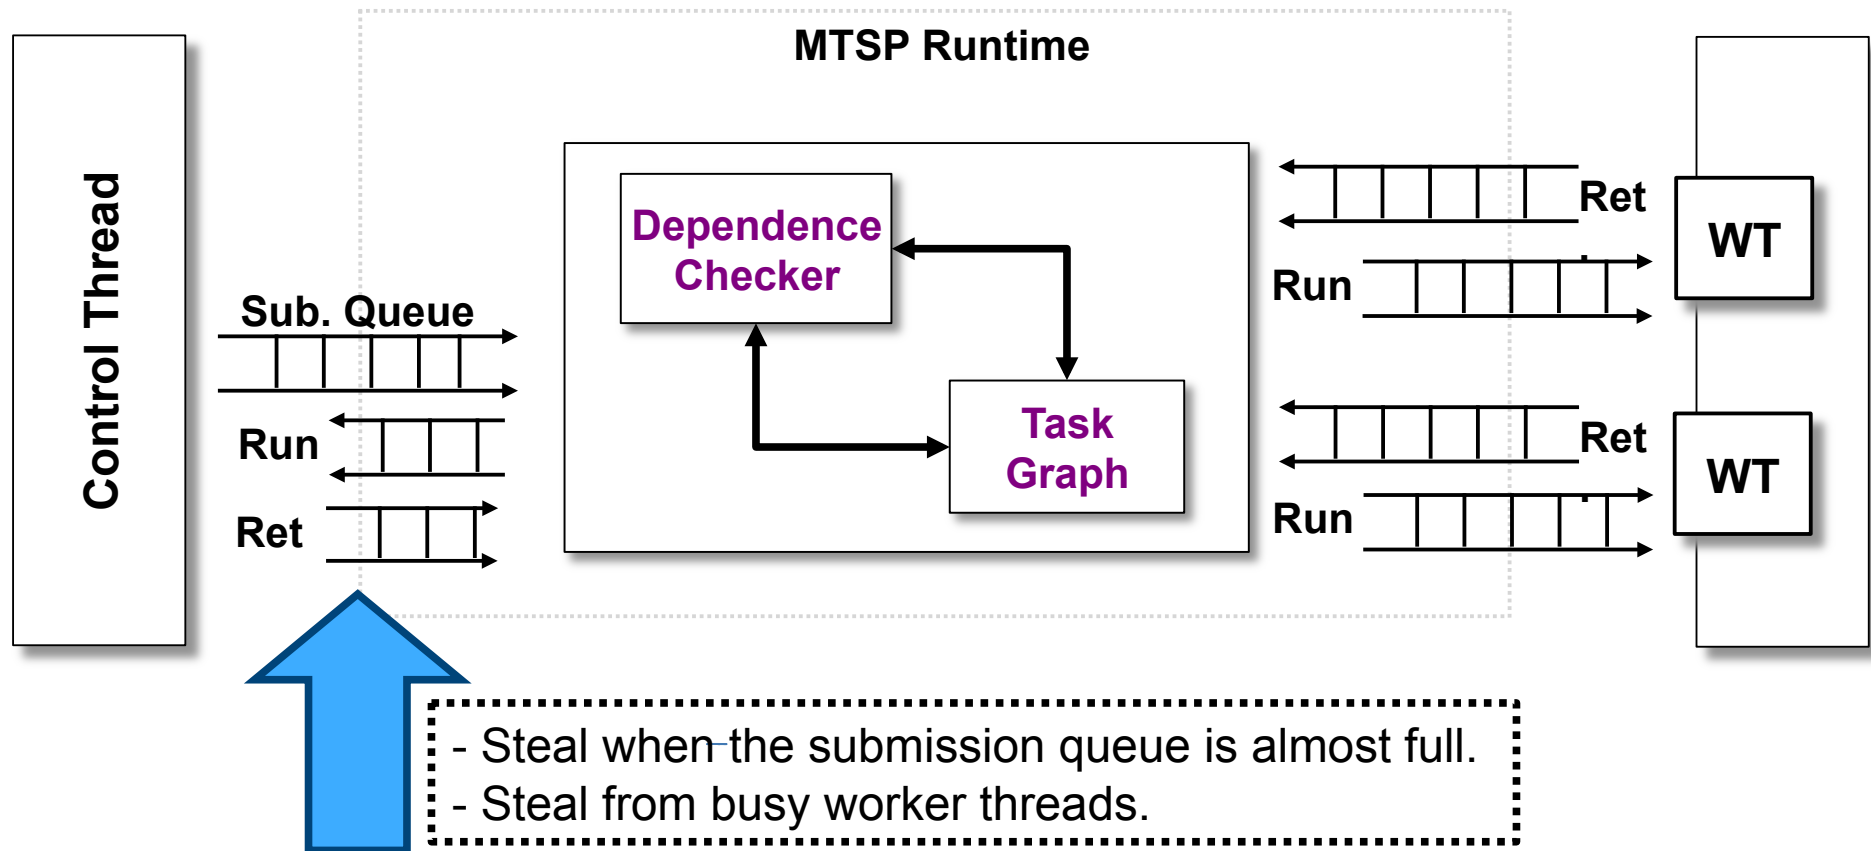


# Control thread waiting


Solution: Work Stealing to control thread



# Work Stealing for the Control Thread



# Control Thread Stealing



**Take away:** for a larger number of threads performance improves!

# Task Parallelism

- Programming model
- Runtime design
- Task stealing
- Runtime performance
- Runtime overhead
- Optimization
  - Task coalescing
  - Dynamic worker deactivation
  - Workload balancing
  - Dependency analysis
- Example

# Efficient Task Execution

# Efficient Task Execution

Manager Cost = 1 second

Task Size = 1.5 seconds

# Efficient Task Execution

Manager Cost = 1 second

Task Size = 1.5 seconds

**1 Worker Thread**

---


# Efficient Task Execution

Manager Cost = 1 second

Task Size = 1.5 seconds

## 1 Worker Thread

---






# Efficient Task Execution

Manager Cost = 1 second

Task Size = 1.5 seconds

## 1 Worker Thread

---



One task is dispatched every 1 second.

One task is consumed every 1.5 seconds.

**No Worker Thread is Idle**

# Efficient Task Execution

Manager Cost = 1 second

Task Size = 1.5 seconds

**2 Worker Threads**

---

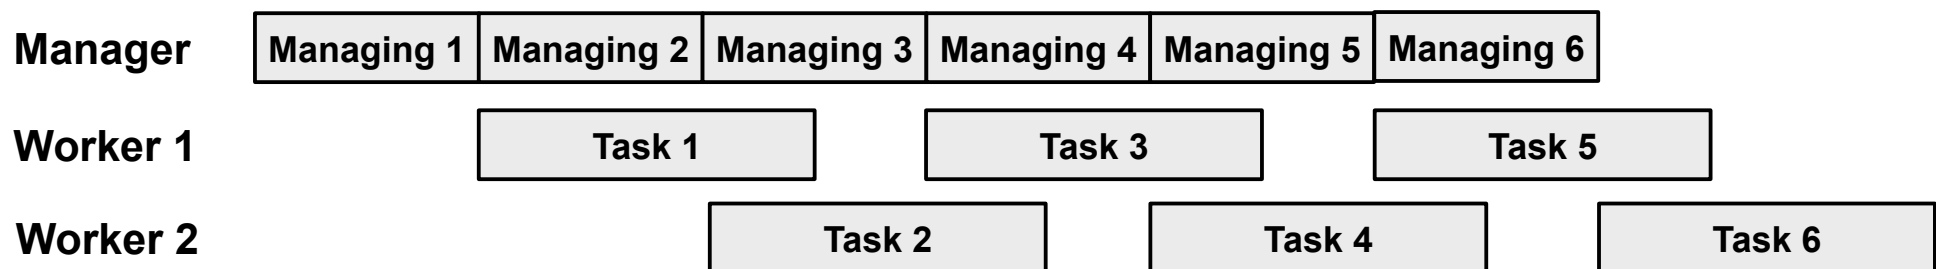
# Efficient Task Execution

Manager Cost = 1 second

Task Size = 1.5 seconds

## 2 Worker Threads

---



**Two tasks are dispatched every 2 seconds.**

**One task is consumed every 1.0 seconds.**

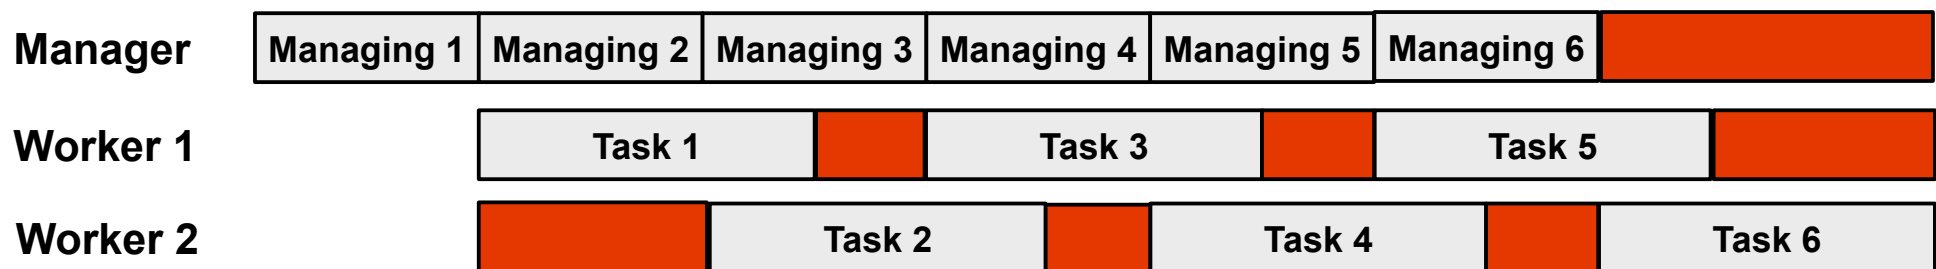
# Efficient Task Execution

Manager Cost = 1 second

Task Size = 1.5 seconds

## 2 Worker Threads

---



**Two tasks are dispatched every 2 seconds.**

**Two tasks are consumed every 1.5 seconds.**

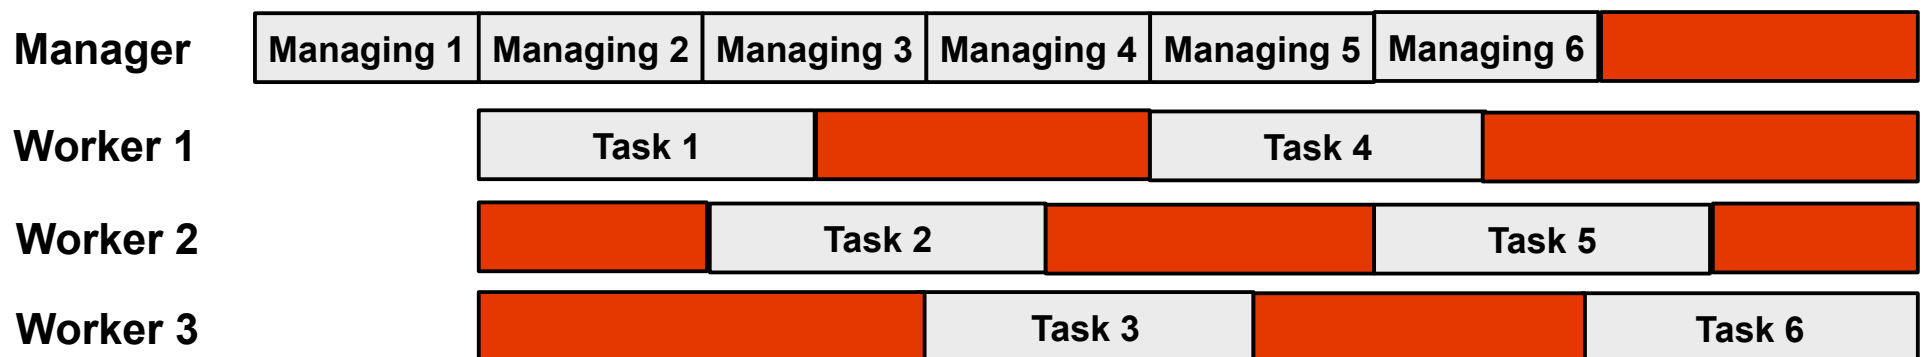
# Efficient Task Execution

Manager Cost = 1 second

Task Size = 1.5 seconds

## 3 Worker Threads

---



# Efficient Task Execution

Manager Cost = **0.5 second**

Task Size = 1.5 seconds


# Efficient Task Execution

Manager Cost = **0.5 second**

Task Size = 1.5 seconds

## 3 Worker Threads

---



# Efficient Task Execution

Rationale:



Task management throughput remained constant.

Task consumption throughput increased


Should activate only the needed number of workers



# The Implementation

- Task Size  TSize
- Runtime Cost  RCost
- Max Number of Workers  ?

# The Implementation

- Task Size  TSize
- Runtime Cost  RCost
- Max Number of Workers  ?

The maximum number of workers for a given *TSize* and *RCost* is:

$$\text{MaxWorkers} = \frac{\text{TSize}}{\text{RCost}}$$

That is, the maximum number of tasks the runtime can dispatch during the execution of another task.

# Task Parallelism

- Programming model
- Runtime design
- Task stealing
- Runtime performance
- Runtime overhead
- Optimization
  - Task coalescing
  - Dynamic worker deactivation
  - Workload balancing
  - Dependency analysis
- Example

# Runtime overhead

Task Data	SparseLU (-64)			
Task ID	402830	402860	4028a0	4028e0
Task cycles	246287	429212	382085	649154
Runtime cycles	9896	19457	25522	33167
#Tasks	64	1024	1024	21856

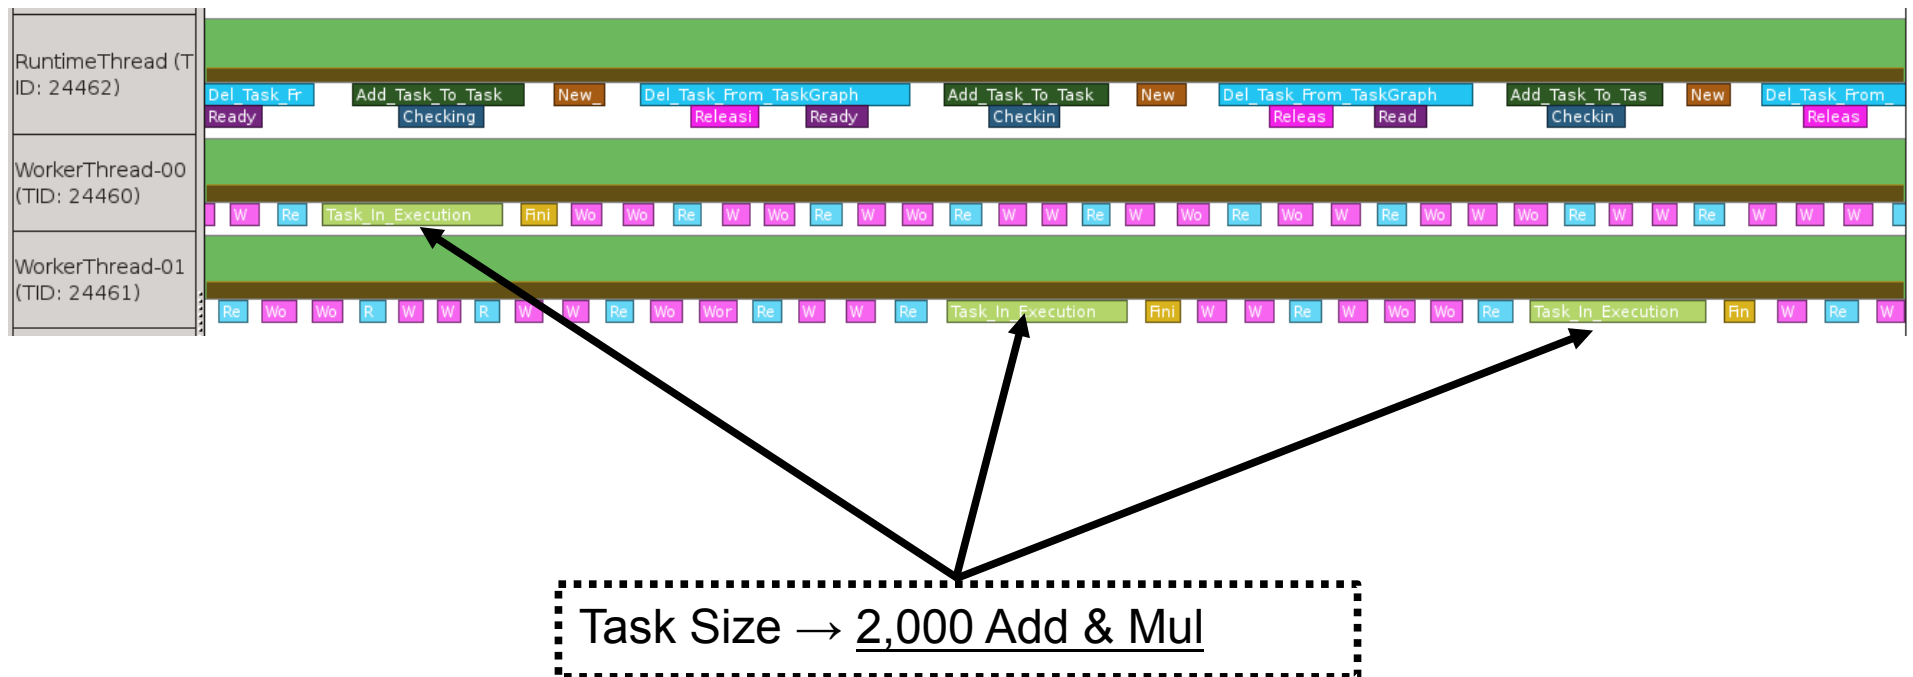
Large tasks

Task Data	Jacobi (-32)			
Task ID	401ab0	401ea0	403040	403170
Task cycles	35736	100728	1682	6742
Runtime cycles	497	637	15234	38907
#Tasks	512	1024	4096	4096

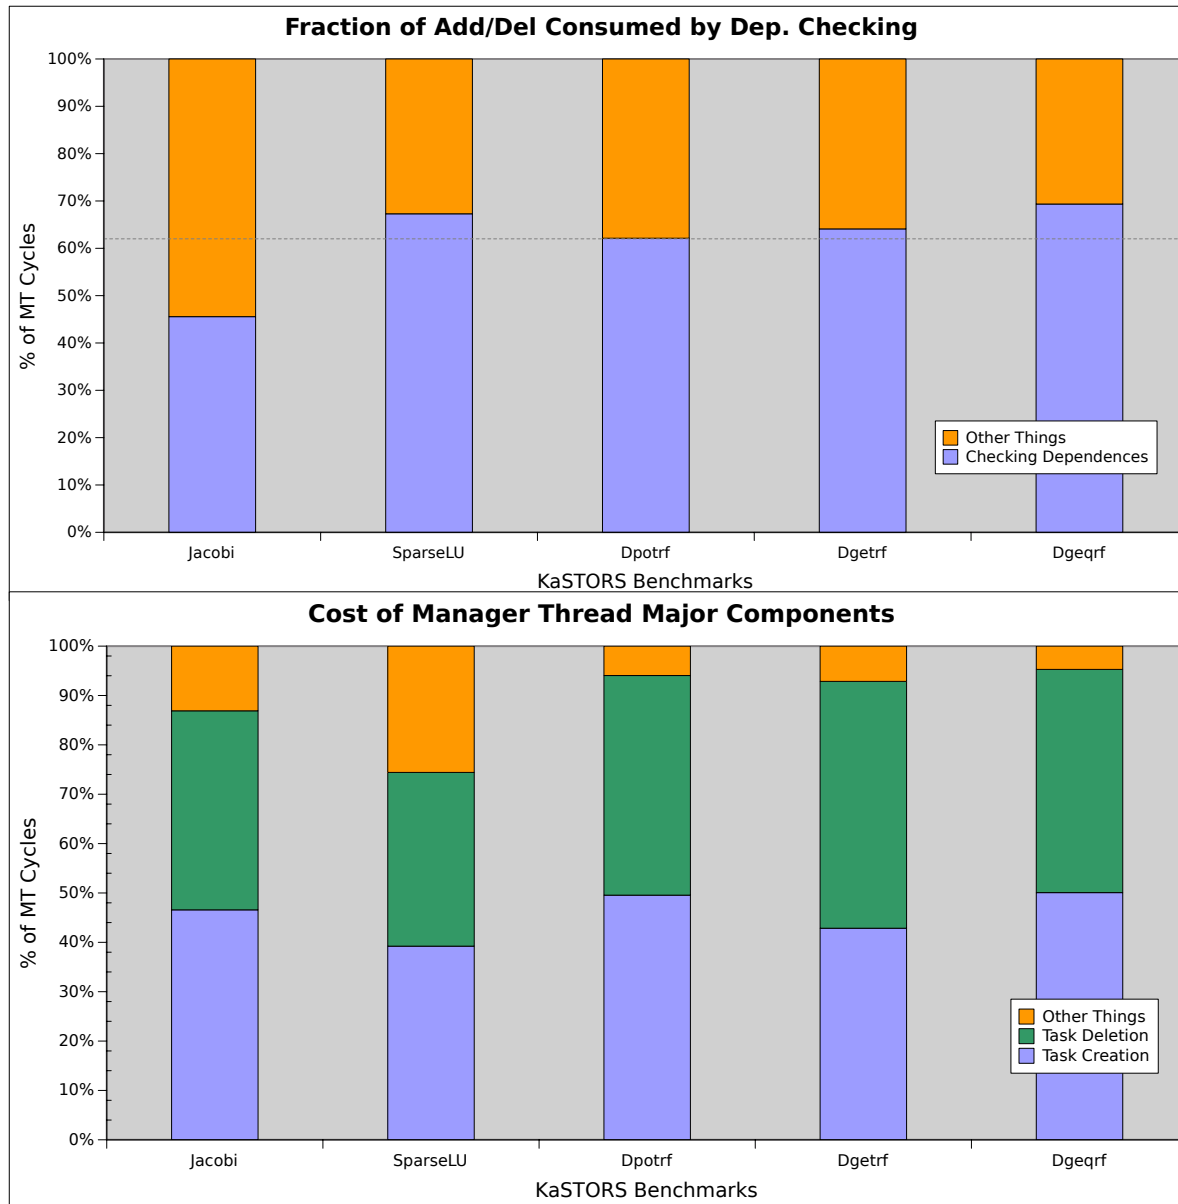
Small tasks with many dependences

# Profiling the MTSP Runtime

# What are the major sources of overhead?



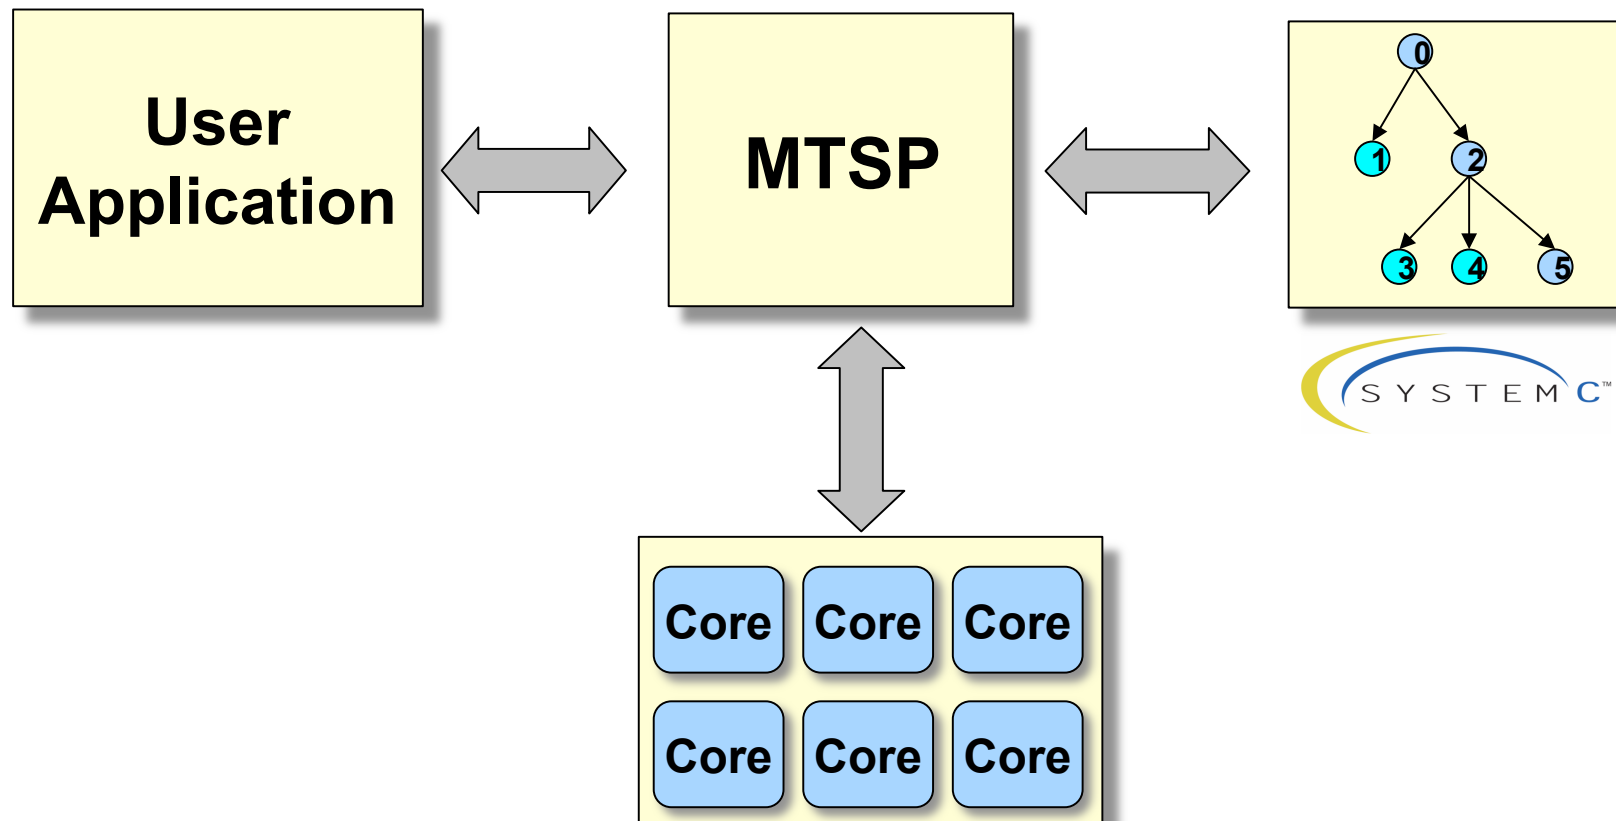
# MTSP Overhead Split



Need to reduce the runtime overhead!

# What if we use hardware support?

Solution: Hardware-enabled task management  
Task Graph Accelerator (TGA)



# Task Parallelism

- Programming model
- Runtime design
- Task stealing
- Runtime performance
- Runtime overhead
- Research
  - Task coalescing
  - Dynamic worker deactivation
  - Workload balancing
  - Dependency analysis
- Example



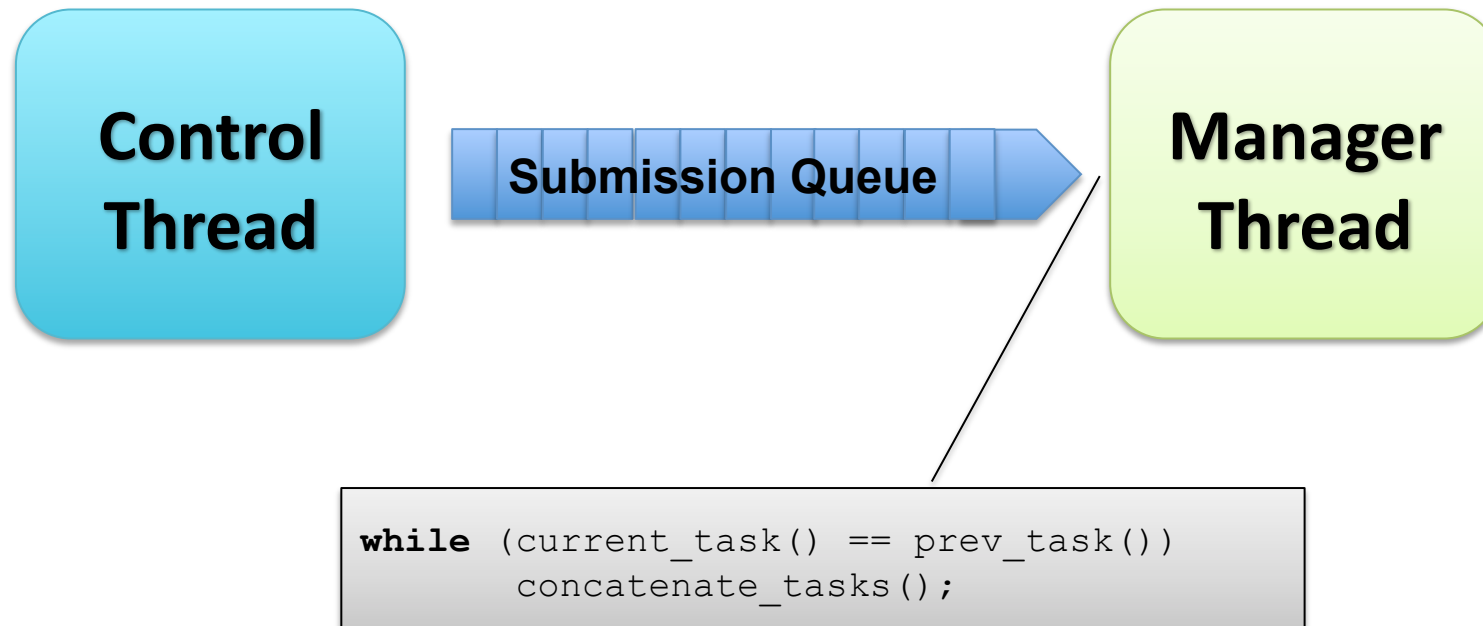
# Task Coalescing

- What is it?
  - Group together subsequent tasks from the submission queue.
- Why do it?
  - Reduce the average time spent with task management.

# Task Coalescing




# Task Coalescing



# Task Coalescing

## Manager Thread

---



Individual task management.

# Task Coalescing

## Manager Thread

---




1 Second...

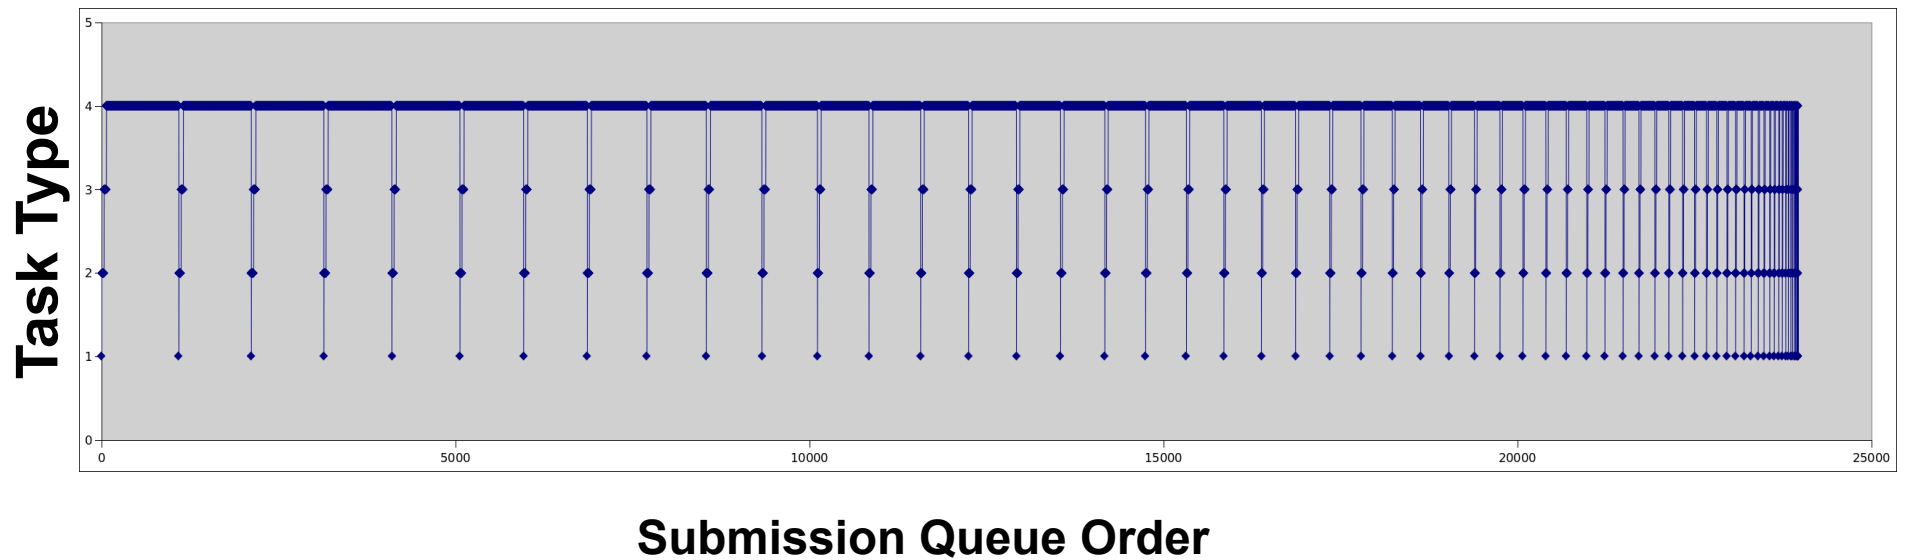
but 2 Tasks

Coalesced task management.


# Task Submission Pattern (Jacobi)



# Task Submission Pattern (SparseLU)




# Parameter Sharing Among Tasks





# Parameter Sharing Among Tasks



# Task Parallelism

- Programming model
- Runtime design
- Task stealing
- Runtime performance
- Runtime overhead
- Optimization
  - Task coalescing
  - **Dynamic worker deactivation**
  - Workload balancing
  - Dependency analysis
- Example

# The Idea

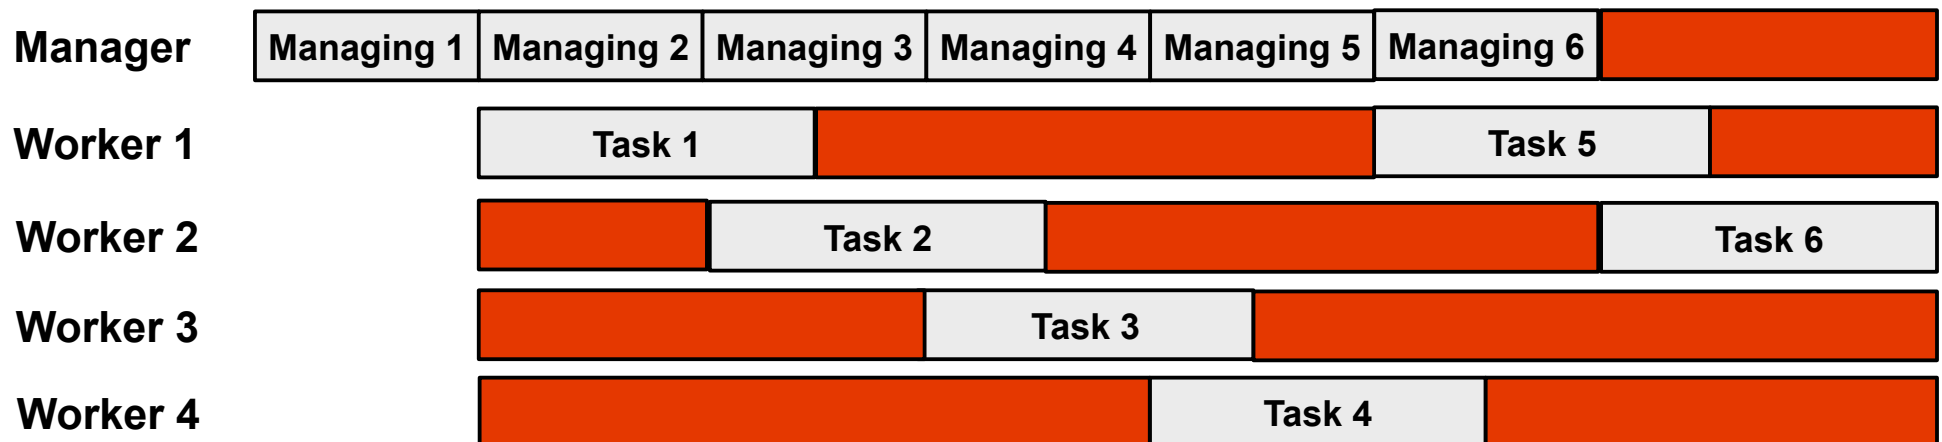
Manager Cost      = 1 second      ← **Difficult to improve!**

Task Size            = 1.5 seconds      ← **Cannot control.**

# The Idea

Manager Cost = 1 second ← Difficult to improve!

Task Size = 1.5 seconds ← Cannot control.

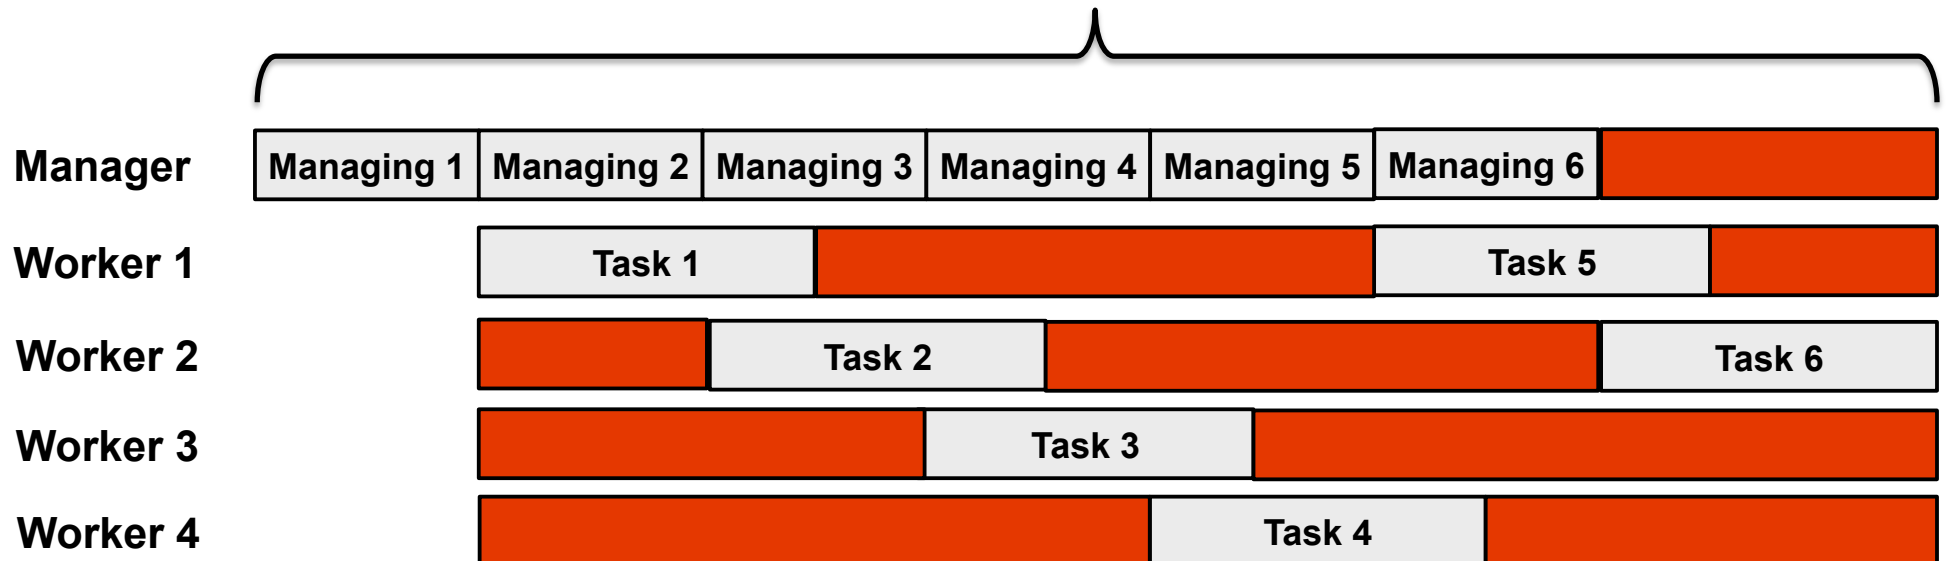


# The Idea

Manager Cost = 1 second ← Difficult to improve!

Task Size = 1.5 seconds ← Cannot control.

## Wall Time




# The Idea

Manager Cost = 1 second ← Difficult to improve!

Task Size = 1.5 seconds ← Cannot control.

$$\text{CPU Time} \approx \text{Wall Time} * 5$$




# The Idea

Manager Cost = 1 second ← Difficult to improve!

Task Size = 1.5 seconds ← Cannot control.

**Can we reduce CPU Time without degrading Wall Time?**



# The Implementation

- Always use a small number of threads!



# The Implementation

- ~~Always use a small number of threads!~~




# The Implementation

- ~~Always use a small number of threads!~~
- Programs have phases...

# The Implementation

- ~~Always use a small number of threads!~~
- Programs have phases...
- Dynamically compute the instantaneous maximum number of workers supported!

# Remembering

- Task Size  TSize
- Runtime Cost  RCost
- Max Number of Workers  ?

The maximum number of workers for a given *TSize* and *RCost* is:


$$\text{MaxWorkers} = \frac{\text{TSize}}{\text{RCost}}$$

That is, the maximum number of tasks the runtime can dispatch during the execution of another task.

# Task Parallelism


- Programming model
- Runtime design
- Task stealing
- Runtime performance
- Runtime overhead
- Optimization
  - Task coalescing
  - Dynamic worker deactivation
  - Workload balancing
  - Dependency analysis
- Example

# Critical Path Prediction and QoS




Which node should execute first? 7, 8 or 9?

# Critical Path Prediction and QoS





# Critical Path Prediction and QoS







# Critical Path Prediction and QoS



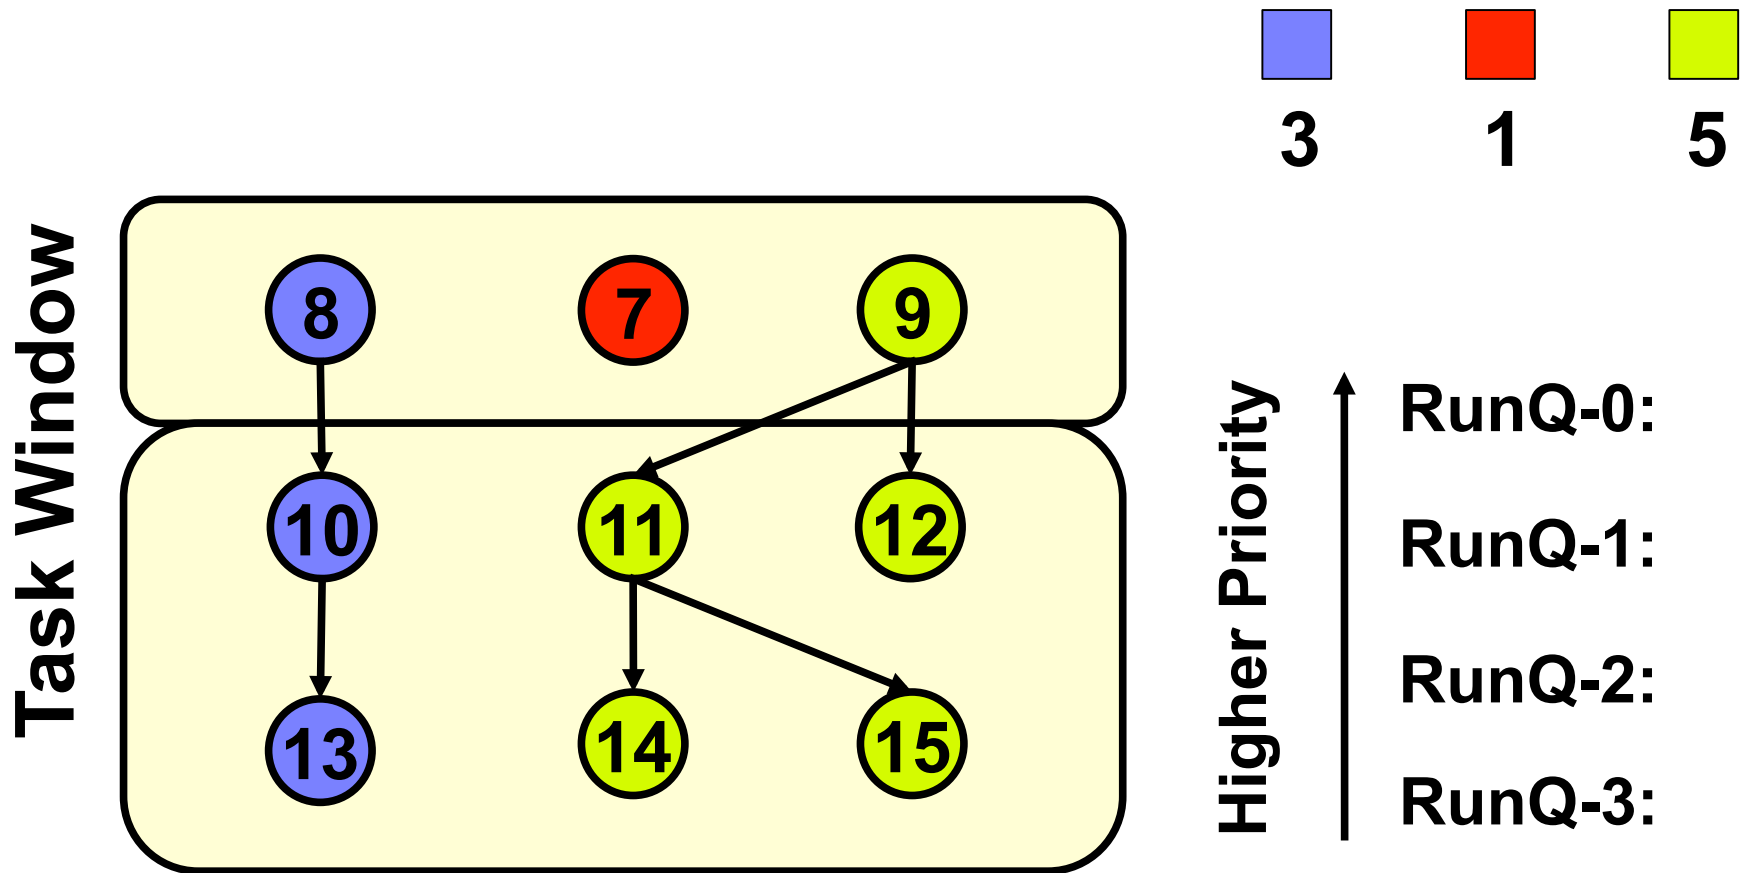
# Critical Path Prediction and QoS



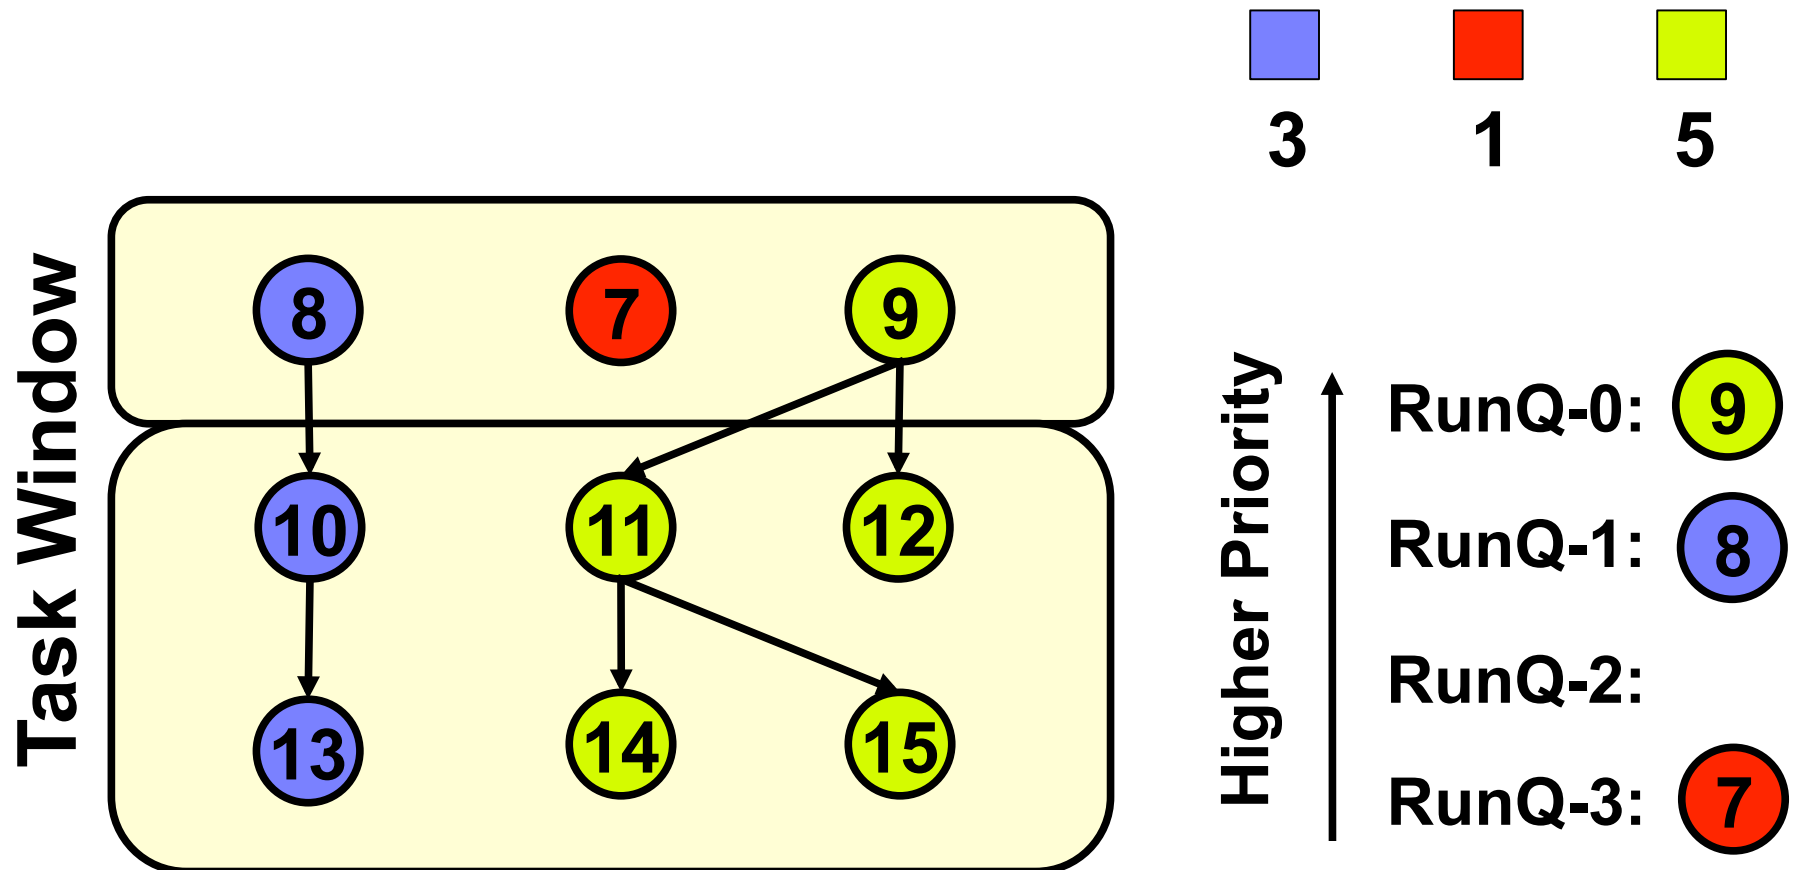
Task Window



# Critical Path Prediction and QoS



# Critical Path Prediction and QoS



# Task Parallelism

- Programming model
- Runtime design
- Task stealing
- Runtime performance
- Runtime overhead
- Optimization
  - Task coalescing
  - Dynamic worker deactivation
  - Workload balancing
  - **Dependency analysis**
- Example

# Programming Support

```
#pragma omp parallel for check
```

```
for (i=0; i < N; i++) {
```

```
    a[i] = fun(&a[i]);
```

```
}
```



What if fun writes  
to a[i+1] ?!

- Programmers make mistakes
- Have to support loop-carried detection
- New check clause to OMP
- IWOMP 2014

# Support Correctness

```
for (i=0; i < N; i++) {  
    #pragma omp task in(u[i]) out(v[i]) check  
    fun(&u[i], &v[i]);  
}
```



What if fun writes  
to u[i] ?!

- We intend to do the same for tasks
- Use escape analysis techniques

# Support Performance

```
for (i=0; i < N; i++) {  
    #pragma omp task in(u[i]) out(v[i]) check  
    fun(&u[i], &v[i]);  
}
```

- Does it pay off?
  - Runtime should estimate the speed-up
  - Evaluate task duration, runtime overhead and # worker threads
  - Develop heuristics to give hints to the programmer
  - It could latter be used for automatic annotation



# Task Parallelism

- Programming model
- Runtime design
- Task stealing
- Runtime performance
- Runtime overhead
- Research
  - Task coalescing
  - Dynamic worker deactivation
  - Dependency analysis
- Example